

Solving the Many-Electron Schrödinger Equation With a Transformer Architecture

Jorge Munoz Laredo

November 14, 2025

Outline

- 1 The Schrödinger Wave Function and the physical laws that rule
 - Schrödinger Equation
 - Physical laws and conditions
 - Optimizing an Ansatz
- 2 Deep Learning Tool Box
 - Natural Gradient Descent
 - Attention Mechanism
- 3 Fermi Net and PsiFormer
 - Fermi Net
 - Psi Former

The Schrödinger equation

On 1926 Schrodinger derived the Time Dependent Form.(TDSE)

$$i\hbar \partial_t \Psi = \hat{H} \Psi,$$

- $\Psi \in \mathcal{H}$ is a complex value function called **wave function**.
- \hat{H} is called the **Hamiltonian Operator**, encodes all the information of the energy of the system.
- Depends on the position \vec{r} of a particle and the temporal evolution (t) .
- Ψ encodes all information about the system; $|\Psi|^2$ gives a probability density that integrates to 1.

Hamiltonian

In the position basis:

$$\hat{H} = \frac{\hat{\vec{P}}^2}{2m} + \hat{V} = -\frac{\hbar^2}{2m} \nabla^2 + \hat{V}$$

Time Dependent Form

When the wave function could be written as:

$$\psi(\vec{r}, t) = R(\vec{r})T(t)$$

TDSE returns you that:

$$T(t) = e^{-iEt/\hbar} \wedge \hat{H}R(\vec{r}) = ER(\vec{r})$$

Where E is the total energy of the system. The eigen-problem becomes obtain R solving:

Find a Ψ , such that:

$$\left[-\frac{\hbar^2}{2m} \nabla^2 + V(\mathbf{r}, t) \right] \psi(\vec{r}) = E\psi(\vec{r})$$

Find the potential V of the system.

Many-Body System

When considering a many-body system, we need to consider the position of each electron like also the spin of it. When considering n bodies, we have:

$$\hat{H}\psi(\mathbf{x}_0, \dots, \mathbf{x}_n) = E\psi(\mathbf{x}_1, \dots, \mathbf{x}_n)$$

With $\mathbf{x}_i = \{\mathbf{r}_i, \sigma\}$, where $\mathbf{r}_i \in \mathbb{R}^3$ is the position of each particle and $\sigma \in \{\uparrow, \downarrow\}$ is the so called spin.

Considerations

- Each particle interact with all the another particles in specific ways.
- For atoms, consider all the protons, electrons and neutrons.
- Solution obey physical laws.

Setting up the Hamiltonian

The first step is obtain a practical form of the **Hamiltonian**.

- Kinetic energy: $T = -\frac{1}{2} \sum_{i=1}^N \nabla_i^2$.
- Electron-nuclear attraction: $V_{en} = -\sum_{i,I} \frac{Z_I}{r_{iI}}$.
- Electron-electron repulsion: $V_{ee} = \sum_{i < j} \frac{1}{r_{ij}}$.

$$\hat{H} = -\sum_{i=1}^N \frac{1}{2} \nabla_i^2 - \sum_{I=1}^M \frac{1}{2M_I} \nabla_I^2 - \sum_{i=1}^N \sum_{I=1}^M \frac{Z_I}{|\mathbf{r}_i - \mathbf{R}_I|} \quad (1)$$

$$+ \sum_{1 \leq i < j \leq N} \frac{1}{|\mathbf{r}_i - \mathbf{r}_j|} + \sum_{1 \leq I < J \leq M} \frac{Z_I Z_J}{|\mathbf{R}_I - \mathbf{R}_J|} \quad (2)$$

Born Oppenheimer approximation helps with.

Fermi-Dirac statistics

All the fermions follow the Fermi-Dirac Statistics, this is.

- Electrons are indistinguishable fermions.
- Exchanging two electrons flips the wavefunction's sign:
 $\Psi(\dots i, j \dots) = -\Psi(\dots j, i \dots).$
- Pauli exclusion: no two electrons can occupy the same

Slater Determinant

We can enforce this using a determinant to enforce an antisymmetric Ψ .

$$\begin{vmatrix} \phi_1^k(\mathbf{x}_1) & \dots & \phi_1^k(\mathbf{x}_n) \\ \vdots & & \vdots \\ \phi_n^k(\mathbf{x}_1) & \dots & \phi_n^k(\mathbf{x}_n) \end{vmatrix}$$

Where ϕ are called spin orbitals

Kato cusp conditions, Jastrow Factor

When two electrons

- Coulomb potentials cause a sharp cusp in Ψ when particles coalesce.
- Electron–nucleus cusp: $\frac{\partial \Psi}{\partial r_{iI}} \Big|_{r_{iI}=0} = -Z_I \Psi(0)$.
- Electron–electron cusp: $\frac{\partial \Psi}{\partial r_{ij}} \Big|_{r_{ij}=0} = \frac{1}{2} \Psi(0)$.

Jastrow Factor $\exp(\mathcal{J})$

In this work we are going to use this specific form:

$$\mathcal{J}_\theta(x) = \sum_{i < j; \sigma_i = \sigma_j} -\frac{1}{4} \frac{\alpha_{par}^2}{\alpha_{par} + |\mathbf{r}_i - \mathbf{r}_j|} + \sum_{i, j; \sigma_i \neq \sigma_j} -\frac{1}{2} \frac{\alpha_{anti}^2}{\alpha_{anti} + |\mathbf{r}_i - \mathbf{r}_j|}$$

Loss: Variational Principle

Try an ansatz and optimize it.

Variational principle states that the energy of any ansatz is greater than the truth energy.

$$E[\Psi] = \frac{\langle \Psi | H | \Psi \rangle}{\langle \Psi | \Psi \rangle} \geq E_0$$

Minimizing $E[\Psi]$ drives the ansatz toward the ground state. E_0 is the true ground energy (minimum possible).

$$\mathcal{L}(\Psi_\theta) = \frac{\langle \Psi_\theta | \hat{H} | \Psi_\theta \rangle}{\langle \Psi_\theta | \Psi_\theta | \Psi_\theta | \Psi_\theta \rangle} = \frac{\int d\mathbf{r} \Psi^*(\mathbf{r}) \hat{H} \Psi(\mathbf{r})}{\int d\mathbf{r} \Psi^*(\mathbf{r}) \Psi(\mathbf{r})}$$

Define:

$$p_\theta(x) = \frac{\Psi_\theta^2(x)}{\int dx' \Psi_\theta^2(x')} \wedge E_L(x) = \Psi_\theta^{-1}(x) \hat{H} \Psi_\theta(x)$$

Then:

$$\mathcal{L}_\theta = \mathbb{E}_{x \sim p_\theta}[E_L(x)]$$

Variational Monte Carlo

To evaluate that expectation we use Monte Carlo estimator.

Quantum Monte Carlo

With the samples $\mathbf{R}_1, \dots, \mathbf{R}_M \sim p_\theta(\mathbf{R})$ we can make:

$$\mathcal{L}_\theta = \mathbb{E}_{x \sim p_\theta}[E_L(x)] \approx \frac{1}{M} \sum_{i=1}^m E_L(\mathbf{R}_k)$$

With:

$$E_L(\mathbf{R}_k) = \frac{\hat{H}\psi(\mathbf{R}_k)}{\psi(\mathbf{R}_k)} = -\frac{1}{2} \frac{\nabla^2\psi(\mathbf{R}_k)}{\psi(\mathbf{R}_k)} + V(\mathbf{R}_k)$$

To obtain those samples we use the **Metropolis-Hastings Algorithm** which handles well high dimensions.

Metropolis-Hastings Algorithm

1. Take a initial configuration $\mathbf{X}_0 \in E$ arbitrary:
2. Propose $\mathbf{X}' = \mathbf{X}_0 + \eta$, where $\eta \sim q(\eta)$, q is a probability density on E called **proposal kernel**. symmetric Gaussian
3. Compute the quantity:

$$A(\mathbf{X}_0, \mathbf{X}') = \min \left(1, \frac{\rho(\mathbf{X}')}{\rho(\mathbf{X}_0)} \frac{q(\mathbf{X}' - \mathbf{X}_0)}{q(\mathbf{X}_0 - \mathbf{X}')} \right)$$

Where ρ is the target distribution where we want sample, In the case where q is symmetric, this simplifies to:

$$A(\mathbf{X}_0, \mathbf{X}') = \min \left(1, \frac{\rho(\mathbf{X}')}{\rho(\mathbf{X}_0)} \right)$$

4. Generate a uniform number $U \in [0, 1]$
5. If: $U < A(\mathbf{X}_0 \rightarrow \mathbf{X}'_l)$ then $\mathbf{X}_1 = \mathbf{X}'$, otherwise try another \mathbf{X}' . Accept or decline.

Metropolis-Hastings

6. Repeat until obtain N_{eq} accepted, the change stabilizes (stationary distribution) this phase is called **burn in**.
7. From $\mathbf{X}_{N_{eq}}$ generate M samples until reach the sample $\mathbf{X}_{N_{eq}+M+1}$. In each sample generates $E_L(\mathbf{R}_k)$ then average to obtain $\mathbb{E}(E_L)$ and begin the back propagation step.

Gradients of the Loss

Using calculus you obtain:

$$\nabla_{\theta}\mathcal{L} = 2\mathbb{E}_{x \sim \Psi^2}[(E_L(x) - \mathbb{E}_{x' \sim \Psi^2}[E_L(x')])\nabla \log|\Psi(x)|]$$

This expectation is calculated in the same way.

Outline

- 1 The Schrödinger Wave Function and the physical laws that rule
 - Schrödinger Equation
 - Physical laws and conditions
 - Optimizing an Ansatz
- 2 Deep Learning Tool Box
 - Natural Gradient Descent
 - Attention Mechanism
- 3 Fermi Net and PsiFormer
 - Fermi Net
 - Psi Former

Fisher Information Matrix

Gradient descent assumes that parameters lives on the Euclidian Space \mathbb{R}^n . ($\mathcal{L}(\theta_1, \dots, \mathbf{x})$)

But in the case where parameters lives on a probability distribution $p(x, \theta)$ that assumption don't work to well. A natural **metric** when looking for compare the distributions $p(x, \theta)$ and $p(x, \theta + \delta\theta)$ is the **Fisher information Matrix**.

Fisher Matrix $F(\theta)$

Let $x \sim p(x|\theta)$. Define the score:

$$s_\theta(x) = \nabla_\theta \log p(x|\theta)$$

$$F(\theta) = \mathbb{E}_{x \sim p(\cdot|\theta)} [s_\theta(x)s_\theta(x)^\top]$$

$s \in \mathbb{R}^d$ a column vector. d number of parameters.

Natural Gradient

Natural Gradient

Using the Fisher Metric, the steepest descent direction of the loss is:

$$-F(\theta)^{-1}\nabla_{\theta}\mathcal{L}$$

The updates becomes:

$$\theta_{k+1} = \theta_k - \eta F(\theta_k)^{-1}\delta_{\theta}\mathcal{L}$$

- Compute F , expensive.
- KFAC ameliorate this with two approximations.

Optimizer: KFAC (natural gradient)

The first approximation is: F_{ij} is assumed zero when θ_i and θ_j are on different layers (neural network). So we just care about the same layer ℓ . Calculus say.

$$\mathbb{E}_{p(\mathbf{X})} \left[\frac{\partial \log p(X)}{\partial \text{vec}(\mathbf{W}_\ell)} \frac{\partial \log p(X)}{\partial \text{vec}(\mathbf{W}_\ell)}^\top \right] = \mathbb{E}_{p(\mathbf{X})} [(\mathbf{a}_\ell \otimes \mathbf{e}_\ell)(\mathbf{a}_\ell \otimes \mathbf{e}_\ell)^\top]$$

Where \mathbf{a}_ℓ are the forward activation and \mathbf{e}_ℓ are the backward sensitivities for that layer.

$$\mathbb{E}_{p(\mathbf{X})} [(\mathbf{a}_\ell \otimes \mathbf{e}_\ell)(\mathbf{a}_\ell \otimes \mathbf{e}_\ell)^\top]^{-1} \approx \mathbb{E}_{p(\mathbf{X})} [\mathbf{a}_\ell \mathbf{a}_\ell^\top] \otimes \mathbb{E}_{p(\mathbf{X})} [\mathbf{e}_\ell \mathbf{e}_\ell^\top]^{-1}$$

Attention on the room

Multi Head Attention

$$\mathbf{o}_{t,i} = \sum_{j=1}^t \text{Softmax} \left(\frac{\mathbf{q}_{t,i}^T \mathbf{k}_{j,i}}{\sqrt{d_h}} \right) \mathbf{v}_{j,i}$$

Figure 1: Original Transformer

$$\mathbf{k}_i = \mathbf{W}_k \mathbf{h}_i, \mathbf{q}_i = \mathbf{W}_q \mathbf{h}_i, \mathbf{v}_i = \mathbf{W}_v \mathbf{h}_i$$

$$\mathbf{u}_t = W^O[\mathbf{o}_{t,1}; \mathbf{o}_{t,2}; \dots; \mathbf{o}_{t,n_h}]$$

Outline

- 1 The Schrödinger Wave Function and the physical laws that rule
 - Schrödinger Equation
 - Physical laws and conditions
 - Optimizing an Ansatz
- 2 Deep Learning Tool Box
 - Natural Gradient Descent
 - Attention Mechanism
- 3 Fermi Net and PsiFormer
 - Fermi Net
 - Psi Former

FermiNet baseline

First deep-neural-network wavefunction ansatz (Pfau et al., 2020). Architecture: multiple dense layers with electron-wise feature streams. Outputs single-electron orbitals feeding into a Slater determinant.

Fermi Net Architecture

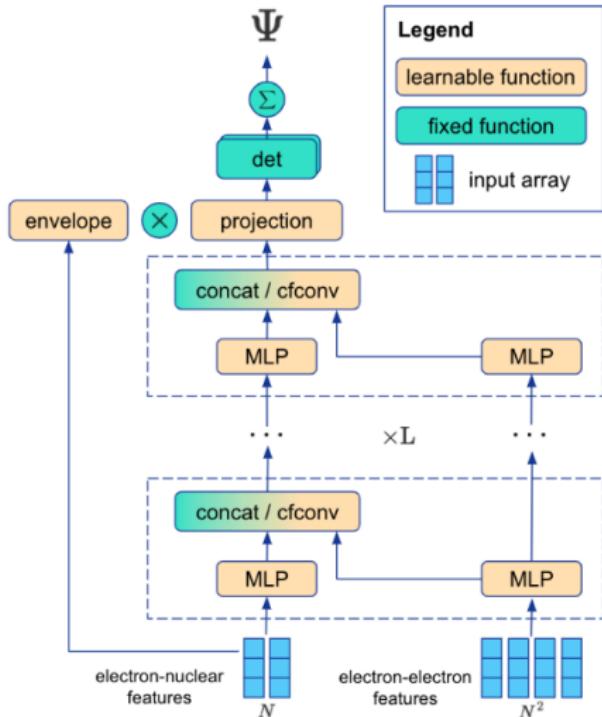


Figure 2: Fermi Net Architecture

The Psiformer Ansatz (overview)

- Neural network trial wavefunction (“Wavefunction Transformer”).
- Uses self-attention layers to model electron correlations.
- Permutation equivariant (independent of electron ordering).
- Greatly improves accuracy, especially for larger systems.
- Solves Schrödinger equation from first principles (no external data).

Psi Former Architecture

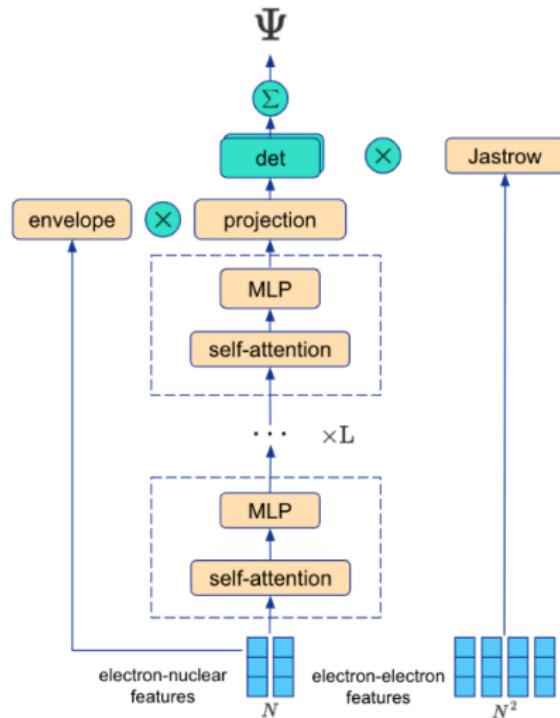


Figure 3: Psi Former Architecture

Psiformer vs FermiNet: Architecture & Attention

- FermiNet layers mix electron features via fixed functions; Psiformer uses self-attention.
- Self-attention: each electron attends to all others (learns interactions).
- Both ansatzes enforce antisymmetry via Slater determinants.
- Psiformer captures correlations more effectively with fewer parameters.

Thanks

Thanks.